



Data-driven versus Topology-driven Irregular Computations on GPUs



Rupesh Nasre

University of Texas
Austin, USA

Martin Burtscher

Texas State University
San Marcos, USA

Keshav Pingali

University of Texas
Austin, USA

What is IrRegularity?

- Data-access or control patterns are unpredictable at compile time.

Irregular data-access

```
int a[N], b[N], c[N];  
readinput(a);  
  
c[5] = b[a[4]];
```


Irregular control-flow

```
int a[N];  
readinput(a);  
  
if (a[4] > 30) {  
    ...  
}
```

Needs dynamic techniques


Pointer-based data structures often contribute to irregularity.

Examples of Irregular Programs




Delaunay Mesh Refinement


```
a = &x  
b = &y  
p = &a  
*p = b  
c = a
```



Points-to Analysis




Single-source shortest paths



Survey Propagation


Operator Formulation




- Computation is modeled as an iterated application of an operator
- **Data-driven processing**: operator is applied only at the nodes where there might be work to be done
- **Topology-driven processing**: operator is applied at all the nodes even if there is no work to do at some nodes

Focus of This Work

- **Study and optimize** topology-driven and data-driven irregular computations on GPUs
 - On GPUs, irregular algorithms are usually implemented in a topology-driven manner
 - On multi-core CPUs, irregular algorithms are usually implemented in a data-driven manner
- **Our findings**
 - Data-driven versions usually perform better
 - Topology-driven versions perform better if additional algorithmic properties can be exploited
 - A hybrid approach can outperform the two



data-driven



topology-driven

Outline

- Motivation
- **Data-driven approach**
- Topology-driven approach
- Hybrid approach
- Evaluation
- Conclusions

Outline

- Motivation
- **Data-driven approach**
 - dual-worklist
 - hierarchical worklists
 - work-chunking
 - atomic-free worklist update
 - work-donation
 - variable kernel configuration
- Topology-driven approach
- Hybrid approach
- Evaluation
- Conclusions


Data-driven: Base Version

```

main {
    read input
    transfer input
    initialize_kernel
    initialize_worklist(wlin)
    clear wlout

    while wlin not empty {
        operator(wlin, wlout, ...)
        transfer wlout size
        clear wlin
        swap(wlin, wlout)
    }
    transfer results
}

```




```


sssp_operator(wlin, wlout, ...) {
    src = wlin[...]
    dsrc = distance[src]
    forall edges (src, dst, wt) {
        ddst = distance[dst]
        altdist = dsrc + wt

        if altdist < ddst
            distance[dst] = altdist
            wlout.push(dst)
    }
}

```




Data-driven: Hierarchical Worklist





- Worklist exploits memory hierarchy
- Makes judicious use of limited on-chip cache

Data-driven: Work Chunking




- Reserves space for multiple work-items in a single atomic
- May reduce overall synchronization

Data-driven: Atomic-free Worklist Update




Data-driven: Work Donation

```
donate_kernel {  
    shared donationbox[...];  
  
    // determine if I should donate  
    --barrier--  
  
    // donate  
    --barrier--  
  
    // operator execution  
  
    // empty donation box  
  
}
```



- Work-donation improves load balance

Data-driven: Variable Kernel Configuration



- Varying configuration improves work-efficiency
- It also reduces conflicts and may improve performance



Outline

- Motivation
- Data-driven approach
- **Topology-driven approach**
 - no global worklists
 - kernel unrolling
 - exploiting memory hierarchy
 - improved memory layout
- Hybrid approach
- Evaluation
- Conclusions

Topology-driven: Base Version

```
main {  
  read input  
  transfer input  
  initialize_kernel  
  do {  
    transfer false to changed  
    operator(...)  
    transfer changed  
  } while changed  
  transfer results  
}
```


cpu gpu



Topology-driven: Kernel Unrolling


```
sssp_operator(src) {  
    dsrc = distance[src]  
  
    forall edges (src, dst, wt) {  
        ddst = distance[dst]  
        altdist = dsrc + wt  
  
        if altdist < ddst  
            distance[dst] = altdist  
    }  
}
```

Memory-bound kernel




- Improves amount of computation per thread invocation
- Need to ensure absence of races
- Propagates information faster

Topology-driven: Exploiting Memory Hierarchy



Topology-driven: Improved Memory Layout




- Bring logically close graph nodes also physically close in memory
- Improves spatial locality

Outline


- Motivation
- Data-driven approach
- Topology-driven approach
- **Hybrid approach**
 - spatial
 - temporal
- Evaluation
- Conclusions

Spatial Hybrid





Shortest paths computation


Spatial Hybrid



Shortest paths computation






data-driven



topology-driven

Temporal Hybrid




Outline

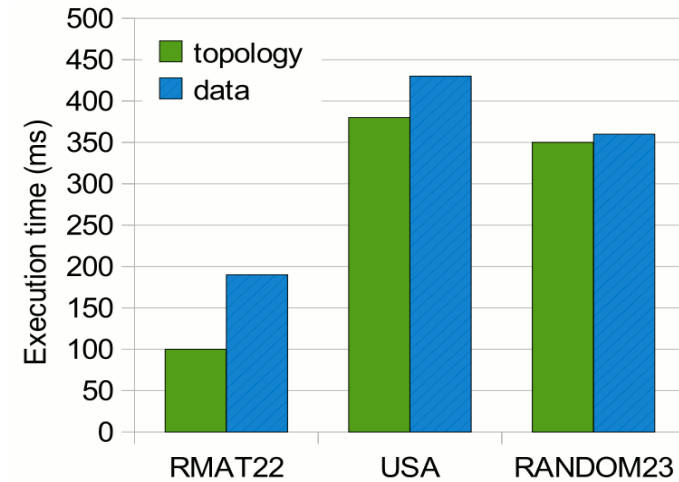

- Motivation
- Data-driven approach
- Topology-driven approach
- Hybrid approach
- **Evaluation**
 - Fermi (Quadro 6000, 1.45 GHz, 448 cores over 14 SMs)
 - **LonestarGPU** benchmark suite
 - <http://iss.ices.utexas.edu/?p=projects/galois/lonestargpu>*
- Conclusions

Data-driven vs. Topology-driven


DMR



SP




BFS




SSSP

Data-driven Optimizations







- Hierarchical worklists, atomic-free worklist update and varying configuration help
- Chunking is not effective
- Overall, performance improves by 40-50%

Topology-driven Optimizations



- Effect is more pronounced than the data-driven approach
- Unrolling helps the most
- Overall, performance improves by 50-60%




Spatial and Temporal Hybrid





- Spatial hybrid has consistently reduced performance
- Temporal hybrid achieves performance better than the solo approaches

Conclusions

- Data-driven irregular computations usually perform better
- Topology-driven irregular computations perform better if additional algorithmic properties can be exploited
- Hybrid approaches may provide better performance than the individual approaches



Data-driven versus Topology-driven Irregular Computations on GPUs



Rupesh Nasre

University of Texas
Austin, USA

Martin Burtscher

Texas State University
San Marcos, USA

Keshav Pingali

University of Texas
Austin, USA

Acknowledgments: Intel, NSF, NVIDIA, Qualcomm

LonestarGPU

ECL
Efficient Computing Laboratory

Galois

DMR Performance

SP Performance

BFS Performance

SSSP Performance